



DATA MINING

Project Report

Bremen Big Data Challenge - Edition 2019

By Gari Ciodaro, Diogo Cosin, and Ralph Florent

May 17, 2019

Abstract

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Laoreet non curabitur gravida arcu ac. Et molestie ac feugiat sed lectus vestibulum mattis ullamcorper velit. Turpis egestas pretium aenean pharetra. Amet justo diam vulputate ut pharetra sit. Elit scelerisque mauris pellentesque pulvinar. Mauris in aliquam sem fringilla. Id ornare arcu odio ut sem nulla. Nunc lobortis mattis aliquam faucibus purus. Eget velit aliquet sagittis id consectetur purus ut faucibus pulvinar. Nulla aliquet enim tortor auctor urna nunc id cursus. Varius quam quisque id diam vel quam elementum pulvinar.

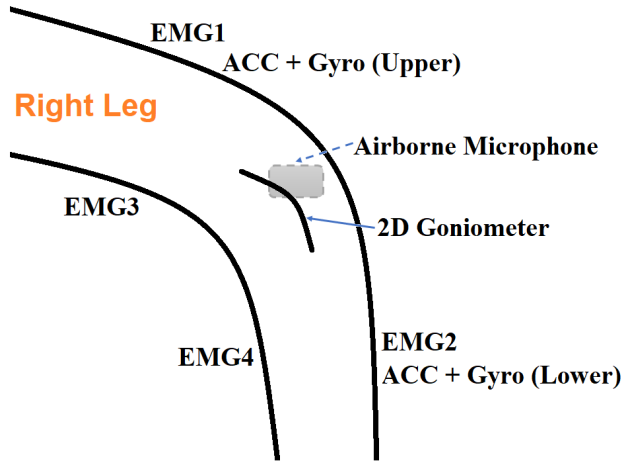


Figure 1: Wearable sensor placement for data measurements. (source: part of the provided data)

1 Introduction

2 Background of the Data

2.1 Data Source

Provided by "The Bremen Big Data Challenge 2019" Organizers, the collected data are based on daily athletic movements [bbdc]. Using wearable sensors above and below the knee (See Figure 1) of the individual (athletic), a dataset of 19 individuals, mainly identified as *subjects*, has been recorded. And as the competition requires, the data of 15 out of the total number of subjects are used as the training dataset and the remaining part as the testing dataset. The dataset is publicly available online on the official website: [BBDC](https://bbdc.csl.uni-bremen.de/index.php/2019h/28-aufgabenstellung-2019) or by simply browsing through the following URL: <https://bbdc.csl.uni-bremen.de/index.php/2019h/28-aufgabenstellung-2019>.

2.2 Description and Format

The data comprise the following 22 movements:

- Race ('run')
- Walking ('walk')
- Standing (standing)
- Sitting ('sit')
- Get up and sit down ('sit-to-stand', 'stand-to-sit')
- Up and down stairs ('stair-up', 'stair-down')
- Jump on one or both legs ('jump-one-leg', 'jump-two-leg')
- Run left or right ('curve-left-step', 'curve-right-step')
- Turn left or right on the spot, left or right foot first ('curve-left-spin-Lfirst', 'curve-left-spin-Rfirst', 'curve-right-spin-Lfirst', 'curve-right-spin-Rfirst')
- Lateral steps to the left or right ('lateral-shuffle-left', 'lateral-shuffle-right')

Subjects	Datafile	Label
Subject02	Subject02/Subject02_Aufnahme000.csv	<i>curve-left-step</i>
Subject02	Subject02/Subject02_Aufnahme001.csv	<i>curve-left-step</i>
Subject02	Subject02/Subject02_Aufnahme002.csv	<i>stand-to-sit</i>
...
Subject19	Subject19/Subject19_Aufnahme438.csv	<i>curve-right-step</i>
Subject19	Subject19/Subject19_Aufnahme439.csv	<i>curve-right-spin-Rfirst</i>

Table 1: Tabular visualization of the "*train.csv*" dataset

Subjects	Datafile	Label
Subject01	Subject01/Subject01_Aufnahme000.csv	<i>X</i>
Subject01	Subject01/Subject01_Aufnahme001.csv	<i>X</i>
Subject01	Subject01/Subject01_Aufnahme002.csv	<i>X</i>
...
Subject15	Subject15/Subject15_Aufnahme438.csv	<i>X</i>
Subject15	Subject15/Subject15_Aufnahme439.csv	<i>X</i>

Table 2: Tabular visualization of the "*challenge.csv*" dataset

- Change of direction when running to the right or left, left or right foot first ('v-cut-left-left', 'v-cut-left-right', 'v-cut-right-left', 'v-cut' right-Rfirst ')

The entire data are available as CSV files, or Comma-Separated Values, and partitioned as training and testing data, respectively represented by the "*train.csv*" and the "*challenge.csv*" files. Starting with the training dataset file (*train.csv*), it contains *UTF-8*¹ character-encoded, line-wise plain texts, whose first line identifies the feature names followed by the feature values. This file contains a total of 6402 lines, which include both the feature names and the feature values. The feature names are *Subjects*, *Datafile*, *Label*, and the feature values map respectively each feature name. For instance, the first feature values of the file are: *Subject02*, *Subject02/Subject02_Aufnahme002.csv*, *stand-to-sit*. Table 1 illustrates a lightweight version of the data partition of the training dataset file.

Similarly, the testing dataset file (*challenge.csv*) is formatted using the same structure with the exception of the *Label* column, which is unknown and marked with an *X*. The datafile contains a total of 1739 lines counting both the feature names and feature values. Table 2 displays a lightweight version of the data partition of the testing dataset file.

Important: Recalling that the dataset is divided into training and testing data, the subjects "*Subject01*, *Subject10*, *Subject14*, *Subject15*" are the selected ones that are used as testing data to assess the solutions. Note the difference in the starting and ending rows of Tables 1 and 2.

As observed in both Tables 1 and 2, each line corresponds to a recording of a movement. The columns have the following meanings:

- **Subject:** The ID of the subject
- **Datafile:** Path of the file containing the sensor data for this recording. For each subject, there is a folder in which individual data files contain the sensor data for individual motion recordings.
- **Label:** The movement that was recorded

Particularly, the *Label* column of the testing dataset contains repeatedly the letter "*X*" to indicate that this value is not present. That is, at the time of submitting solutions, the submission should

¹Unicode Transformation Format, extended ASCII, variable-width encoding.

exactly match the testing data, where each X will be replaced by a label. This label corresponds to the classification result of a specific movement. It is important that the spelling (including upper / lower case) of the textual labels matches exactly the spelling of the labels in the training data.

As mentioned above, the datafiles are references to other CSV files. For example, the path file `Subject02/Subject02_Aufnahme000.csv` is a CSV file within a folder named *Subject02* located in the root path (i.e., the current directory of the downloaded zip files). The CSV file itself is a dataset with a proper format. Basically, the file has a set of comma-separated numbered-values that looks like this: *32688, 32224, 32991, 32609, 32790, 33048, 37168, 34610, 27374, 29068, 29264, 28408, 31784, 28133, 29295, 29244, 33216, 37140, 34736*.

Each line has 19 values (numbers) and represents the sensor values measured at one time (sampled at 1000 Hz). In other words, the columns represent the individual wearable sensors recording the human activities (see Figure 1):

EMG1	Airborne	Goniometer X	Goniometer Y	Gyro lower X
EMG2	ACC upper X	ACC lower X	Gyro upper X	Gyro lower Y
EMG3	ACC upper Y	ACC lower Y	Gyro upper Y	Gyro lower Z
EMG4	ACC upper Z	ACC lower Z	Gyro upper Z	

The size of the CSV datafiles varies inappropriately. That is, in most cases, due to inaccurate measurements, random initialization states, mechanical flaws, computational and processing cost, and so on.

3 General comments

3.1 Notation

Let us defined a modeling procedure as a function $\mathcal{P}(\Theta) : \mathbb{R}^m \longrightarrow \mathbb{R}^k$ where m is the number of features, k is the number of classes, and $\Theta : \{Preprocessing_{method}, Feature_{extration}, Statistical_{technique}\}$ is a set representing parameters. Giving this abstraction Θ modifies the structure \mathcal{P} but always takes a *feature vector* $X \in \mathbb{R}^m$, and returns a *Probability vector* $Y \in \mathbb{R}^k$ where each component $\in [0, 1]$.

It is clear that the $\mathcal{P}(\Theta)$ that represents exactly the reality regarding athletics movements is unknown to us, which let us to defined a measure of the amount of veracity that a giving \mathcal{P} has compare to reality.

$$accuracy = \frac{\sum Correct_{classification}}{N} \quad (1)$$

3.2 Cross validation

Given the modeling procedure \mathcal{P} defined in the Subsection 3.1 for a given learning task, one usually seeks to evaluate its performance in the so-called training data and testing data. The training data is not but the set of data points utilized during the training procedure whereas the testing data is a completely new data points set. Ideally, an efficient modeling procedure \mathcal{P} will fit the training data and also generalize well to new data. However, during training time only training data is available. In this sense, the statistical method *Cross validation* can be applied to estimate training and testing error even though only training data is available. This way, it is possible to estimate if the model is overfitting or underfitting using only the training data set.

Despite being simple, Cross validation is still a powerful method that can be applied to estimate the testing error. The simplicity is in the process. The Training data set is artificially split into two subsets. One partition is then used as the training data while the other one is used as the testing data.

Apart from this simplest data partition strategy, one may also split the data set into K partitions, more commonly called *folds*. The process follows by holding one the of *folds* as the testing data set

while the remaining $K - 1$ folds are aggregated as the training data set. The process is repeated K times given that each partition is utilized as the testing data set exactly once. The model is then assessed by taking the average of the results obtained during the K repetitions. The optimal model will perform well in the testing and training data, avoiding this way overfitting and underfitting, respectively. In more educated terms, it will present low variance when the different K training data sets are introduced by the cross validation, and low bias if it fits the training data set efficiently.

3.3 Curse of dimensionality

The curse of dimensionality is a common machine learning problem. It happens when the dimension of the *input vectors*, m , is much higher than the number of data points N . In this case, the data points will be highly sparsely distributed in the m -dimensional coordinate system called defined by the vector space spanned by the linear combination of the *input vectors*. In other words, the data points will present large distances from each other. Consequently, the modeling procedure \mathcal{P} will have difficulties in finding similarities between the data points.

A solution for this curse is to reduce the dimensionality of the by applying feature extraction functions such as Principal Component Analysis and K -means. The idea in this stage is to reduce the data points dimension without highly compromising in loss of information. A rule of thumb practically applied is to respect a ratio of 10 features per data point.

4 Data Preprocessing

Two *Preprocessing methods* procedures were implemented.

Preprocessing method 1:

1. Take a **subject file**(each file contains a class movement). It can be viewed as $[19 \times N_{fr}]$ matrix composed by 19 columns vectors $S_{data} \in \mathbb{R}^{N_{fr} \times 1}$ where number of records in file is $N_{fr} \in \mathbb{N}^{>0}$.
2. Transpose each S_{data} into a row vector, concatenating them into one single vector $S_{concat} \in \mathbb{R}^{1 \times 19 \times N_{fr}}$
3. Repeat step 1 and 2 for every subject file.
4. Create a dataset with the rows of step 3 with it according it corresponding label (extracted from the name of the file). This data set is a matrix D of dimensions $[6401 \times (19 * N_{fr})]$. For $N_{fr} = 56810$ D is $[6401 \times 1079390 + 1]$.

Preprocessing method 2:

1. Take a **subject file**(each file contains a class movement). It can be viewed as $[19 \times N_{fr}]$ matrix composed by 19 columns vectors $S_{data} \in \mathbb{R}^{N_{fr} \times 1}$ where number of records in file is $N_{fr} \in \mathbb{N}^{>0}$.
2. For each S_{data} decomposed it as $s_{data} = \mu + \omega$ where μ is a smoothed version of s_{data} calculated using lowess with tree points average weighted linear regression. A complete derivation of this algorithm can be found in [1]. Having s_{data} and μ calculate $\omega = s_{data} - \mu$. Create a vector $\mu_{sta} \in \mathbb{R}^3$ with first tree statistical moments of μ , that is, the average, the variance, and the steepness. calculate the discrete Fourier transformation of ω , extract the first five coefficients and put them in a vector $\omega_{fft} \in \mathbb{R}^5$. A complete derivation of this algorithms can be found in [2]. take μ_{sta} and concatenate it with ω_{fft} into a row vector $S \in \mathbb{R}^8$
3. concatenate each S into a single vector $S_{row} \in \mathbb{R}^{8 \times 19}$
4. Repeat steps 1 to 3 for every **subject file** and stack the vectors S with its corresponding label into a matrix D is $[6401 \times 8 * 19]$.

For *Preprocessing method 1* the parameter N_{fr} had to be set, since each **subject file** had different number records. By observing a 100 random sample of files, we concluded rather arbitrarily that $N_{fr} = 56810$ was reasonable number of records. In case a particular file did not meet this requirement, the signal per sensor would repeat itself until the desired N_{fr} was reached. The idea with *Preprocessing method 2* was to capture the general properties of the movement (using μ_{sta}) in terms of its statistical moments. On the other hand, we also attempted to capture information about the periodicity of the movement using ω_{fft} .

5 Data Exploitation

Depending on the *Preprocessing method* some *Feature extraction* and *Statistical Technique* combination would more adequate to implement than other. Here we only present highest *accuracy* combinations.

5.1 Data Exploitation with *Preprocessing method 1*

One problem that is immediately observed with *Preprocessing method 1* is the immense dimensionality of the sample space, therefore trying to find a *Statistical Technique* capable of learning an adequate decision function is basically impossible. To bypass this problem implemented a *Feature extraction* and *Statistical Technique* shown in the Figure 2.

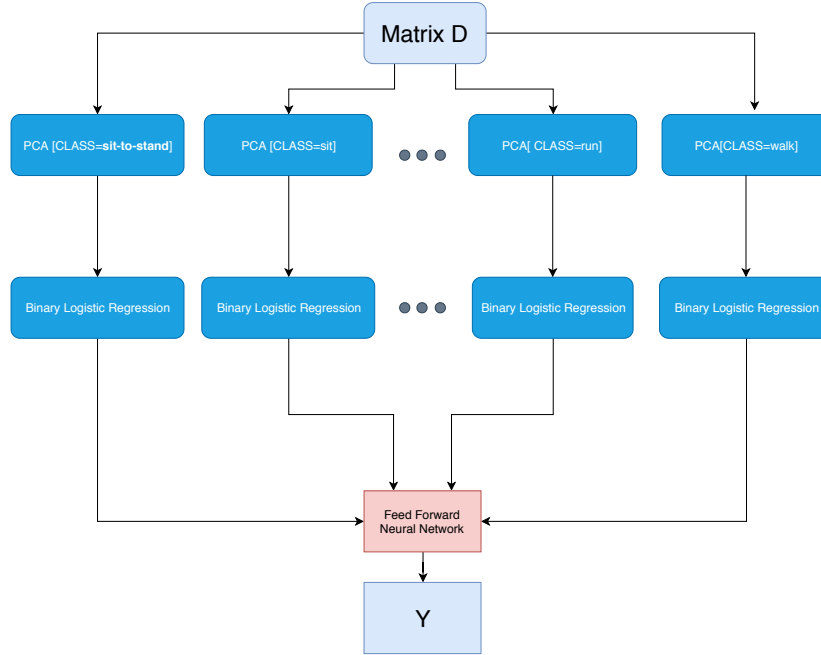


Figure 2: $\mathcal{P}(\Theta_1)$ Scheme

We theorized that the importance of features in matrix D measured by its variance will strongly depend on the particular movement involved, acknowledging this we filtered matrix D by class, and applied principal component analysis $PCA_{|class} : \mathbb{R}^{1079390} \rightarrow \mathbb{R}^{275}$ to extract the first 275 principal components (this transformation will be the beginning a branch in Figure 2) that accounts to 90 % of the variance. Note that the first level in 2 has 22 $PCA_{|class}$ extractors. For training, matrix D is passed without filtering per label to each $PCA_{|class}$ and then feed to a *Binary logistic regression* $LogR_{|class} : \mathbb{R}^{275} \rightarrow [0, 1]$ where the classes are encoded in *one vs all* manner. Each of the $LogR_{|class}$ will output a degree of believe that a particular X belongs to a class k . Finally concatenating the




Figure 3: curve-left-spin-Lfirst




Figure 4: curve-left-spin-Rfirst




Figure 5: curve-left-step




Figure 6: curRight-spin-Lfirst




Figure 7: curRight-spin-Rfirst




Figure 8: curRight-step




Figure 9: jump-one-leg




Figure 10: jump-two-leg




Figure 11: lateral-shuffle-left




Figure 12: lateral-shuffle-right




Figure 13: lay




Figure 14: run

degree of believe predicted by each $LogR_{|class}$ we construct a $X' \in \mathbb{R}^{22}$. Finally X' is passed to a single feed forward neuronal network $NN : \mathbb{R}^{22} \rightarrow \mathbb{R}^{22}$ that predicts our final probability vector Y in a *hot encoded* structure, we take the maximum component and assign its corresponding label to that register.

$$\Theta_1 := \{Preprocessing_{method\ 1}, PCA_{|class}, LogR_{|class}, NN\} \quad (2)$$

5.2 Data Exploitation with $Preprocessing_{method\ 2}$

The dimension of X using $Preprocessing_{method\ 2}$ is R^{152} so no further dimension reduction is required. We implemented a variety of feed forward neuron networks using tensor flow, where we varied the Network architecture composed by the number of hidden units $H_{units} \in \mathbb{N}^{>0}$ per hidden layer $H_L \in \mathbb{N}^{>0}$ the regularization parameter $\alpha \in \mathbb{R}^{>0}$, and the activation functions, namely $\{Logistic, tanh, relu\}$. From Figures 3 to 25 we can observe μ signal examples per label.

$$\Theta_2 := \{Preprocessing_{method\ 2}, NN\} \quad (3)$$




Figure 15: sit




Figure 16: sit-to-stand




Figure 17: stair-down




Figure 18: stair-up




Figure 19: stand




Figure 20: stand-to-sit




Figure 21: v-cut-left-Lfirst




Figure 22: v-cut-left-Rfirst




Figure 23: v-cut-right-Lfirst




Figure 24: v-cut-right-Rfirst




Figure 25: walk

6 Data Analysis

6.1 Data Analysis with $\mathcal{P}(\Theta_1)$

The last procedure to complete the full scheme in Figure 2 is the implementation of the *Fully-Connected Neural Network (FCNN)*. Applying this *Deep Learning* approach on the preprocessed data serves as an alternate technique to optimize the loss function and increase accuracy.

When implementing FCNN using *Scikit-Learn*, there are several specifications to account for. Among them figure the *activation function*, the regularization factor *alpha*, and the number of hidden layers. To briefly summarize some parameters, as per relevance or rule of thumbs, we describe their optional values:

- The solver for the weight optimization ('lbfgs', 'sgd', 'adam')
- The number of hidden layers ranging between 100 ± 20
- The regularization factor varying between 0.1 and 0.00001²
- The activation function for the hidden layers ('identity', 'logistic', 'tanh', 'relu')

Given the preprocessed data, whose order of magnitude scales to 1000+ data points, the default parameters reveal themselves ideal to test out first. Then, based on the obtained results, we would assess the accuracy of the prediction by tuning the other parameters (e.g. alpha) accordingly with the expectations of optimizing the classifier and improving the accuracy.

The settings for the default values when running the MLPClassifier are:

Solver	Activation Function	Hidden Layers	Alpha
<i>lbfgs</i>	<i>relu</i>	<i>50</i>	<i>0.01</i>

Surprisingly, both the training and testing errors score an error rate of 83.7%. Hoping to find the parameters that work best for the task optimization, we performed *Grid Search* by combining above-mentioned options and tweaking them in the same range as in our first test. For instance, the chosen values for alpha were 10^{-1} , 10^{-2} , 10^{-3} , 10^{-4} , and 10^{-5} .

Finally, the best parameters set found on the classifier development set are:

Solver	Activation Function	Hidden Layers	Alpha
<i>lbfgs</i>	<i>tanh</i>	<i>20</i>	<i>0.01</i>

with an improved accuracy of 84.44%. This leads to believe that there is a non-linear relationship between the feature vectors and the classified labels.

6.2 Data Analysis with $\mathcal{P}(\Theta_2)$

Giving that the same *Statistical Technique*, namely *Fully-Connected Neural Network (FCNN)* was used for $\mathcal{P}(\Theta_2)$ the prior section parameters definitions holds the same however, instead of using *Scikit-Learn* we set up an environment where the *FCNN* algorithm was implemented using *keras* with *Tensor-flow* back-end for *gpu* usage.

6.2.1 experimental setup

Train a *FCNN* with such high dimension dataset D [6401×153] is computationally expensive (in fact this is the reason we were forced to use *keras*) so trying a *cross-validation* plus *grid-search* as in $\mathcal{P}(\Theta_1)$ is not practically possible. We used instead a simply 75% and 25% train-test split, and randomly tested different *FCNN* architectures. We also min-max scale so that *FCNN* was trained using features $\in [0, 1]$ interval, for increasing the convergence time.

²Practice shows that by varying alpha the classifier has greater chance to perform way better than by adjusting the number of layers, or trying other algorithm.

7 Results and Discussions

7.1 Results and Discussions for $\mathcal{P}(\Theta_2)$

$H_{L1}(H_{units}, fun)$	H_{L2}	H_{L2}	α	$Train_{Score}$	$Test_{Score}$	$Challenge_{score}$
60,relu	87,tanh	23, sigmoid	0.000001	95%	91%	64%

Table 3: Results on $\mathcal{P}(\Theta_2)$




Figure 26: $\mathcal{P}(\Theta_2)$ Confusion matrix on test data.

From Figure 26 we can conclude that the most misclassification occurs with the classes *stand* and *sit* which intuitively are mechanically similar, this can be further confirmed by looking at the smoothen (μ) plots 15 and 19. We tried to modified the *Preprocessing method*₁ to include some manually crafted features that will indicate a clear different between *stand* and *sit*, they can be summarized as follows:

- The first and second derivate of μ , since we believed the direction of the movement had to be different.
- The maximum value of the μ , since sensor data related to distance from the floor must different for both classes.
- increase to 5 five the number of statistical moments.

Unfortunately none of them gave better results on the test data. The main reason being that this particular adjustment had to me made to all the features, this in term caused other features to get noise, as a result, the decrease in misclassification gained in the pair *stand, sit* was undermined by the increase of misclassification on other features, particularly *sit-to-stand* and *stand-to-sit*.

7.2 Results and Discussions for $\mathcal{P}(\Theta_1)$

In Table 4 we can observe that the Challenge score is inferior to that reached on $\mathcal{P}(\Theta_2)$, since this was the case, we focused our efforts of optimization on $\mathcal{P}(\Theta_2)$.

H_{units}	Activation function	α	CV_{train}	CV_{test}	Challenge score
20	tanh	0.01	(??)	84.44%	54%

Table 4: Results on $\mathcal{P}(\Theta_1)$

8 Conclusion

8.1 Conclusion for $\mathcal{P}(\Theta_1)$

- In retrospective we realized that we introduced some *data leakage* on this implementation, since while constructing $PCA_{|class}$ we feed the algorithms the whole dataset. This could be an explanation for the difference in performance on our CV_{test} and challenge score.
- In *Preprocessing method 1* we arbitrarily choose the number of time steps to be 56810 per signal, we believed that this decision should be made with more empirical knowledge of the phenomena, to increase the performance of $\mathcal{P}(\Theta_1)$.
- We also implemented *Boosting* with *random forest* with 4% increase in CV_{test} however, the competition had already finished and we could not test the result on the challenge data.

8.2 Conclusion for $\mathcal{P}(\Theta_2)$

- After the competition finished we realized that one reason for the difference in performance between our test and the challenge score might have been that the splitting of our data was done randomly, we believe that setting aside entire subjects would have been a better approach, after all, an individual should walk and run in a similar manner.
- In other to increase the performance of $\mathcal{P}(\Theta_2)$ we believe a window of sampling the signal S_{data} should be implemented using expert knowledge on the subject. As a side note, we tried the latest by splitting S_{data} into 3 sectors and then repeating *Preprocessing method 2*. unfortunately, this is very computational expensive, and the attempt, after days of running had to be stooped.

References

- [1] Cleveland, W.S. (1979) “Robust Locally Weighted Regression and Smoothing Scatterplots”. *Journal of the American Statistical Association* 74 (368): 829-836.
- [2] Oraintara, S., Chen, Y. J., & Nguyen, T. Q. (2002). Integer fast Fourier transform. *IEEE Transactions on Signal Processing*, 50(3), 607-618.